
000
001
002
003
004
005
006
007
008
009
010
011
012

Dragon-Alpha&cu32: Fully Open-source Java Deep-Learning Framework And High-Performance CUDA Library

013
014
015
016
017
018
019
020
021
022
023
024
025
026
027
028
029
030
031
032
033
034
035
036
037
038
039
040
041
042
043
044
045
046
047
048
049
050
051
052
053
054

Abstract

Java is powerful, but its potential in deep-learning (DL) realm remains untapped. At present, Python DL frameworks are predominantly preferred over those based on Java, mainly attributed to their superior usability and flexibility. To leverage Java's capabilities, Dragon-Alpha, a Java tensor computing framework, has been developed. The primary goal of Alpha is to provide ease of use, high-performance, and scalability. Alpha offers high-level user-friendly APIs for DL, and other levels of APIs to meet diverse requirements. Alpha achieves GPU acceleration via cu32 library, and improves efficiency through techniques like async APIs, C-K-S algorithm, Im2col-Winograd, and fused operators. Alpha has the potential to consolidate computing power across heterogeneous platforms and devices, based on its multi-layer architecture and compatibility with Java big-data ecosystem. Alpha&cu32 are fully open-source, with no reliance on external libraries. Experiments show that Alpha&cu32 surpass PyTorch&cuDNN on Cifar10 and ILSVRC2012 in certain scenarios.

1. Introduction

Deep learning (DL) is a highly hot field of artificial intelligence. Deep neural networks (DNNs) have been widely applied in diverse fields, yielding remarkable achievements. However, the impressive capabilities of DNNs come with increased complexity. To simplify the representation of DNNs and accelerate their execution, DL frameworks have been developed.

Python DL frameworks ([Paszke et al., 2019](#); [Abadi et al., 2016](#); [Jia et al., 2014](#); [Tokui et al., 2015](#); [Chollet](#)) have gained immense popularity, due to their user-friendly and

flexible nature. On the other hand, Java DL frameworks ([Gibson et al., 2016](#); [Feenster et al.](#)) are not as popular as their Python counterparts, mainly because of their more intricate APIs and limited adaptability. However, Java itself has many benefits, including robustness, speed, flexibility, platform independence, community support, and a strong big-data ecosystem, suggesting that Java's potential in DL field has not been fully exploited.

For efficiency, DL frameworks integrate acceleration libraries (aclibs) to take advantage of parallel processors like GPUs, CPUs, and FPGAs. Nevertheless, many widely-used GPU aclibs such as cuBLAS and cuDNN ([Chetlur et al., 2014](#)) are not open-source. The details of high-performance GPU programming remain somewhat opaque. As a result, there needs open-source GPU aclibs, that can help users better understand and utilize GPUs.

To address these issues, Dragon-Alpha&cu32 have been developed. As a Java tensor computing framework, Alpha can be used to express and execute DL algorithms. Cu32 is an efficient aclib for *float32* computations on GPUs, and has been integrated into Alpha. Alpha&cu32 are fully open-source, and only require JDK and CUDA for execution. Besides, cu32 adopts C-K-S algorithm and Im2col-Winograd to accelerate convolution.

In experiments conducted on Cifar10 and ILSVRC2012 datasets, Alpha&Cu32 exhibit comparable efficacy and convergence to PyTorch&cuDNN, and outperform them in certain circumstances.

2. Background

2.1. Deep Learning Frameworks

DNNs are multi-layer neural networks, with complex computational graphs and numerous parameters. Their complexity allows the learning and expression of sophisticated patterns, but brings challenges for concise representation and rapid execution. DL frameworks are designed to tackle these problems, with ease of use and efficiency as their key priorities.

DL frameworks offer a set of reusable primitives, based on the commonly used structures of DNNs. These primitives can be linked to aclibs for faster processing. Users are

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

055 required to define the forward propagation of DNNs using
 056 these primitives, while the automatic differentiation takes
 057 care of the backpropagation. Such propagation can be
 058 conceptualized as dataflows. Many DL frameworks, like
 059 Caffe (Jia et al., 2014), TensorFlow (Abadi et al., 2016),
 060 Keras (Chollet), and DeepLearning4j (DL4j) (Gibson et al.,
 061 2016), use static flows, whereas PyTorch (Paszke et al.,
 062 2019) and Chainer (Tokui et al., 2015) use dynamic flows.
 063 Static flows are pre-compiled and remain constant
 064 throughout the repeated batch calculations, while dynamic
 065 flows are created by the interaction between layers in
 066 forward propagation. Static flows are more predictable,
 067 making them easier to optimize and potentially leading to
 068 higher performance; dynamic dataflows have greater
 069 flexibility and usability, as they have larger state space.
 070 The encapsulation of aclibs provides convenience, but may
 071 hide certain specifics and restrict functionalities. While
 072 operators in DL frameworks are easy to use, it may be
 073 challenging to fine-tune computing strategies, use operators
 074 beyond the parameter constraints, or avoid redundant
 075 mechanisms in particular scenarios.

2.2. Java or Python

076 At present, the primary advantages of Python DL
 077 frameworks over the Java ones are user-friendliness and
 078 flexibility. Java frameworks, like DL4j and Deep Java
 079 Library (DJL) (Feenster et al.), typically adhere to
 080 conventional Java programming patterns, which are
 081 proficient in managing complex relations, but may be
 082 cumbersome for progress-oriented tasks such as executing
 083 DNNs. Relatively, their APIs are more complex and overly
 084 encapsulated. In DL4j and DJL, DNNs are created using a
 085 series of chained invocations to builders, allowing for
 086 flexible hyperparameter configuration but may need a
 087 substantial amount of code. The training process is
 088 encapsulated into objects, including Trainer and Listener,
 089 which promotes standardization but may conceal certain
 090 details and limit customized options. In DL4j, the model,
 091 optimizer, and loss function are configured in a single
 092 builder, instead of completely decoupled, potentially
 093 compromising flexibility.

094 Conversely, Java frameworks have advantages over the
 095 Python ones. They can be integrated into Java big-data
 096 ecosystem to harness computing power across different
 097 platforms, particularly beneficial for training large models.
 098 Besides, they leverage Java's advantages over Python's
 099 limitations. Java is a compiled language with static
 100 datatype, while Python is an interpreted language using
 101 dynamic datatype. Given their nature, Java is generally
 102 more efficient, robust, and secure. Java uses multiple
 103 threads for parallelism, whereas Python relies on costly
 104 processes. Therefore, Java potentially has higher
 105 parallelism and better hardware utilization. At times,
 106 Python's simplicity may obscure specifics and impose
 107

108 constraints, while Java's complexity can offer advanced
 109 capabilities for tailoring solutions to diverse cases.

110 By capitalizing on Java's strengths and adopting innovative
 111 designs, it's feasible to develop a DL framework with ease
 112 of use and efficiency.

3. Characteristics

To develop Alpha&Cu32, the goal is to attain a balance of usability, high-performance, and scalability. The focus is not only on the advancement of specific technologies, but also on the overall coherence of the entire system.

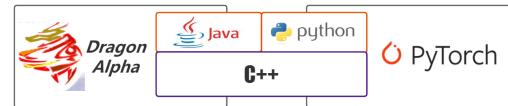


Figure 1. Alpha is implemented by C++ at bottom and Java at top

3.1. Different Levels of APIs

Alpha supports different levels of APIs, including user-friendly high-level APIs and complex low-level APIs. The top-level APIs enable Alpha to work as a DL framework; the lower-level APIs serve as intermediary zones between the top-level APIs and aclibs, offering better control over system specifics. Users can overcome the limitations of higher-level APIs via lower-level APIs, to directly manipulate memory addresses, adjust computing strategies, and circumvent certain mechanisms. Based on different levels of APIs, it's feasible to develop various applications, even a different DL framework.

3.2. Easy-to-use

Alpha's high-level APIs are succinct and easy to use, incorporating the facade pattern, factory pattern, and chained invocation. To facilitate programming and identifying key information, the naming conventions are brief and minimize the use of capital letters. The coding style combines the elements of Java and Python, making it recognizable to users of both languages.

Alpha supports dynamic dataflows for DNNs. The provided operators resemble mathematical formulas in their structure. Hyperparameters can be specified using constructors and factories, and further modified through chain invocations. To construct a directed computational graph of DNNs, users can easily create a subclass of Module and override the *forward* method. Backpropagation can be accomplished with the assistance of automatic differentiation. Figure 2 illustrates a relevant example.

3.3. Scalability

As a Java application, Alpha is platform-independent and can be integrated into Java big-data frameworks such as Hadoop (Cutting et al.) and Spark (Paszke et al., 2010). Its multi-layer architecture separates the hardware specifics from the top layers, enabling polymorphism for different

```

public static class Block extends Module {
    Unit conv1 = new Unit("conv1", conv3D(1, 3, 3, 1, 1));
    public void forward_in_channel(int in_channel, int out_channel, int stride) {
        conv1 = conv3D(true, in_channel, out_channel, 3, stride, 1);
        bnl = batchNorm(out_channel);
        if(stride > 1 || out_channel != in_channel)
            downsample = nn.Sequential(
                conv3D(false, in_channel, out_channel, 3, stride, 1),
                batchNorm(out_channel));
    }
}

@Override
public void __init__(Engine eg) {
    super.__init__(eg);
    for(BatchNorm bn : this.find(BatchNorm.class))
        bn.affine(true).beta(0.1f).beta2(0.1f).eps(1e-5f);
}

@Override
public Tensor[] __forward__(Tensor... x) {
    Tensor[] res = x;
    X = F.leakyRelu(bnl.forward(conv1.forward(x)));
    if(downsample != null) res = downsample.forward(res);
    return F.leakyRelu(bn.add(X, res[0]));
}
}

public static class Network extends Module {
    Unit conv1 = nn.conv3D(3, 64, 3, 1, 1);
    Unit bnl = nn.batchNorm(false, 64);
    Block block1 = new Block(64, 128, 2);
    Block block2 = new Block(128, 256, 2);
    Unit fc = nn.fullconnect(true, 256, 10);
}

@Override
public Tensor[] __forward__(Tensor... x) {
    X = F.leakyRelu(bnl.forward(conv1.forward(x)));
    X = block1.forward(X);
    X = block2.forward(X);
    X = F.adaptive_avgPool2D(1, X);
    return fc.forward(F.flatten(X));
}

static { alpha.home("alpha-home"); }
static Mempool memp = alpha.engine.semopl(alpha.MEM_1GB * 8);
static Engine eg = alpha.engine.cuda_float32(0, memp, alpha.MEM_1MB * 2048);

Network net = new Network();
net.train().init(eg, println);
Optimizer opt = alpha.optim.Adam(net.param_map(), 1f, println);
LossFunction ls = alpha.loss.softmax_crossEntropy();
BufferedTensorIter iter = Cifar10.train().buffered_iter(eg, batchsize);

alpha.stat.load_zip(net, "weight");
alpha.stat.opt.load_zip(opt, "opt_weight");
eg.sync(false);
for(int i=0; i<epoch; i++)
    for(iter.shuffle_sort().reset(); iter.hasNext(); ) {
        PairTensor pair = iter.next();
        Tensor x = pair.input;
        Tensor y = pair.label;

        Tensor yh = net.forward(x)[0];
        alpha.print("loss = ", ls.loss(yh, y));
        net.backward(ls.gradient(yh, y));
        opt.update(0).clear_grads();
        net.gc();
    }
}

alpha.stat.save_zip(net, "weight");
alpha.stat.save_zip(opt, "opt_weight");

```

Figure 2. Using Alpha’s high-level APIs to construct and train DNNs

types of devices. Alpha can utilize specific devices based on the appropriate underlying layers, while maintaining the stability of upper layers.

3.4. High-performance

Alpha utilizes GPUs through cu32, whose kernel functions are well-optimized. Some kernels are general solutions to ensure minimum performance thresholds, while others are designed to enhance effectiveness in specific contexts.

Alpha's async APIs enable the concurrent execution of operators, leading to enhanced parallelism and hardware utilization. Alpha can optimize static dataflows within dynamic dataflows, through inplace operators, fused operators, parallel execution of branches, and etc.

Alpha uses memory-pools to cache memory blocks, to minimize the overhead required for memory-allocation system-calls. Memory-pools work on the abstract *malloc* and *free* methods, which are implemented by lower-level components. A certain amount of pinned memory is managed to act as a cache between JVM and devices, to expedite data transmission through direct memory access. To save bandwidth, images are stored&transferred using *int8*, and converted to float datatypes at destination devices.

4. Architectures

As shown in Figure 3, Alpha is composed of 7 layers.

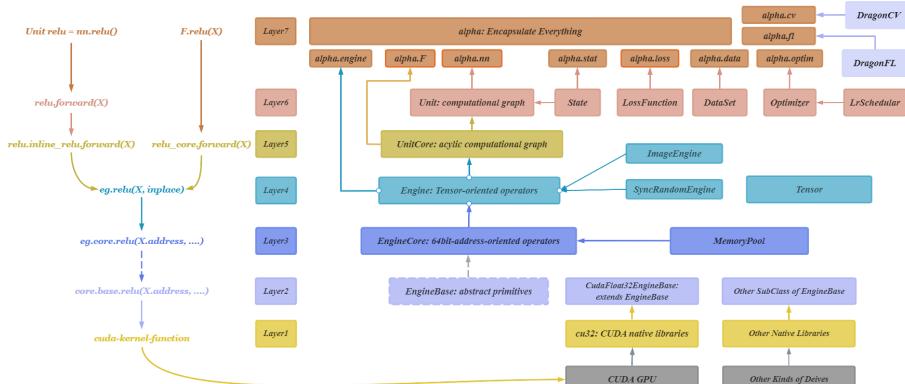


Figure 3. Alpha's 7-layer architecture

165 error-handling, where the memory-pool is abstract and
 166 customizable. EngineCore can work on specific devices via
 167 suitable EngineBases. For instance, It can perform *float32*
 168 operations on GPUs using *CudaFloat32EngineBase*.
 169

4.4. Engine

171 Engine encapsulates EngineCore to enhance usability, by
 172 warping parameters and 64bit addresses in Tensor objects.
 173 The last dimension of Tensors is implicitly padded to 4x,
 174 to vectorize memory access using 128bit units. Tensors
 175 play a crucial role in automatic differentiation, and act as
 176 semaphores to coordinate operators. Tensors can be released
 177 manually or automatically by JVM (Java virtual machine).
 178 'Sync' and 'check' properties act as the switches of async
 179 APIs and parameter-check. Typically, the variability of
 180 computational graphs in DNNs is limited and predictable.
 181 Therefore, checking parameters at the initial few batches is
 182 necessary; at every time could be wasteful. Once certain
 183 batches have been processed to ensure correctness, such 2
 184 properties can be set to false, to enable async APIs and
 185 disable parameter-check, thereby improving efficiency.
 186 ImageEngine provides operators for image processing and
 187 augmentation, with the ability to deal with hyper-spectrum
 188 images. SyncRandomEngine generates random numbers
 189 using 1 thread, whereas Engine uses multi-threads. These 2
 190 methods may result in different quality of random numbers.
 191

4.5. UnitCore

192 UnitCores are applications of Engines, which are used to
 193 construct the acyclic computational graphs of DNNs. Their
 194 public *forward* method establishes connections between
 195 UnitCores via callbacks, to represent the directed edges.
 196 The protected *backward* method is invoked by automatic
 197 differentiation. In this process, UnitCores gather gradients
 198 from their successors along the resolved edges, and then
 199 calculate gradients for their predecessors.
 200

4.6. Unit

201 Units are built upon UnitCores, which have weights and
 202 hierarchical structures. A cyclic graphs formed by Units
 203 can be transformed to an acyclic graph composed of
 204 UnitCores. Compared to the garbage collection of JVM, the
 205 *gc* method of Unit is more cost-effective and timely for
 206 recycling resources. Units collaborate with DataSets,
 207 LossFunctions, and Optimizers to train DNNs. DataSet
 208 supports multi-thread data loading and preprocessing,
 209 which can run in parallel with DNN-execution through a
 210 buffer. DragonCV and DragonFL are developed for basic
 211 image and file processing.
 212

4.7. Alpha

213 *alpha* packaged almost everything. The functionalities can
 214 easily accessed via keywords like *alpha.engine*, *alpha.nn*,
 215 *alpha.F*, *alpha.data*, *alpha.optim*, *alpha.loss*.
 216

5. Techniques

5.1. Asynchronous APIs

Alpha's asynchronous (async) APIs enable the parallel
 execution of multi operators. For example, one CPU thread
 can launch many CUDA kernels without blocking. When
 executing operators on devices, JVM can undertake work
 such as generating logs and managing data structures.
 Async APIs are easy to use. Users only need to focus on the
 relation among operators, rather than underlying details.

When async APIs are activated, operators promptly return
 result Tensors, regardless of whether the computation has
 been completed. Each Tensor is associated with a Syncer,
 that can be triggered by *Tensor.c* method. Syncers act as
 semaphores to wait for the end of corresponding operators,
 and can be programmed to perform tasks like resource
 reclamation. Async APIs can be used to optimize dataflows,
 by processing unrelated multi branches simultaneously. An
 illustrative example is shown in Figure 4.

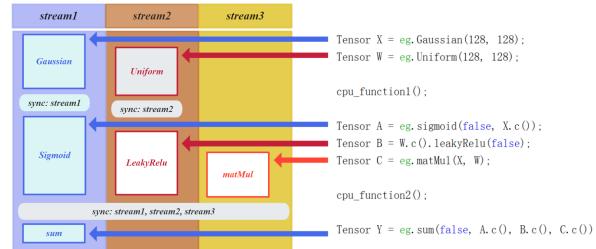


Figure 4. Async APIs enable the concurrent execution of multi operators alongside CPU instructions. At most 3 operators can be executed simultaneously in 3 CUDA streams. 2 CPU functions are performed in parallel with the GPU.

5.2. C-K-S Algorithms: skip zeros in convolution

Zero items (0s) are commonly included in convolutional
 operators, such as 0-padding on feature-maps, and
 0-insertion in deconvolution and dilated convolution. These
 0s cause unnecessary calculations and strain on hardware.
 To skip these 0s, C-K-S algorithm is proposed, including
 ConvV2, KS-deconv, and Sk-dilated. ConvV2 excludes
 padded 0s, providing a constant factor speedup; KS-deconv
 and SK-dilated transform sparse tensors to dense ones,
 accelerating ND deconvolution and dilated convolution by
 $stride^N$ and $dilate^N$ times. C-K-S is based on math and
 not reliant on specific systems. Its operations have minimal
 interdependence to comply with the nature of SIMD.

KS-deconv is designed for non-unit-stride deconvolution, to
 bypass ($stride - 1$) 0s inserted between adjacent items of
 input-feature-maps (*ifms*). As shown in Figure 5,
 KS-deconv has 3 stages: the filters are split into $stride^N$
 parts; each part is used to conduct stride-1 convolution on
 the corresponding subset of *ifms*; the resulting outputs are
 composed to obtain the output-feature-maps (*ofms*). This
 kernel-split idea is based on 2 observations. Firstly, among
 the patches of *ifms*, the 0-distributions can be categorized

220
221
222
223
224
225
226
227
228
229
230
231

Figure 5. KS-deconv in backpropagation of 2D conv-layers. In stage1, $W(3 \times 3)$ is rotated 180 degrees, and split to $C_{00}(2 \times 2)$, $C_{01}(2 \times 1)$, $C_{10}(1 \times 2)$ and $C_{11}(1 \times 1)$. C_{00-11} are concatenated to $C(2 \times 2 \times 2 \times 2)$ which has continuous memory. In stage2, C_{00-11} are respectively used to perform stride-1 convolution with ∇Y , and the outputs are O_{00} , O_{01} , O_{10} and O_{11} . In stage3, O_{00-11} are composed to obtain ∇X .

into $stride^N$ classes, which can be distinguished by the (*coordinates* $\%stride$) of these patches. Secondly, each item in *ofms* can be calculated, by performing a dot-product on two nonzero segmentations where are respectively derived from the filters and *ifms*.

Sk-dilated is used for dilated convolution, where sparse filters have ($dilate - 1$) 0s placed between adjacent items. In each filter, except for the *channel* and *batch* axes, the coordinates of nonzero items must be integral multiples of *dilate*. Adhering to this rule, Sk-dilated does not add 0s to filters. Instead, for each dot-product of dilated convolution, it fetches items in filters with unit step-size, and selects items in *ifms* with steps-size of *dilate*.

ConvV2 utilizes filter-trimming to exclude padded 0s at the boundary of *ifms*, thereby enabling access only to nonzero elements in the central region. This technique is achieved by moving pointers and constraining memory access, without auxiliary workspace. Filter-trimming’s efficacy is positively correlated to the size of 0-padding, so it usually plays a greater role on small feature-maps.

Filter-trimming has been integrated in KS-deconv and Sk-dilated, to develop their V2 version. As illustrated in Figure 6, KS-deconv-V2 are compared with PyTorch-1.12 in backpropagation on RTX 3060ti. As the feature-maps become smaller, the speed of KS-deconv-V2 increases.

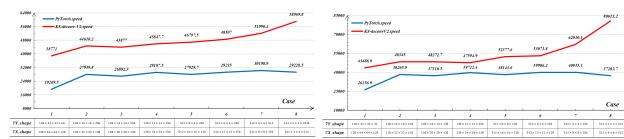
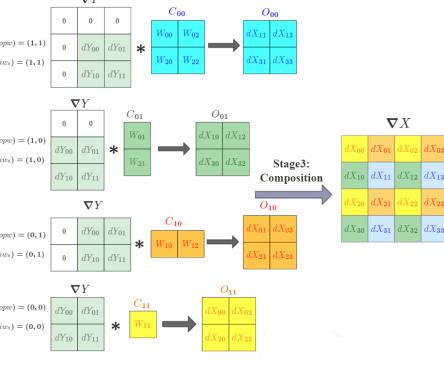


Figure 6. Compare KS-deconv-V2 with PyTorch. The stride is 2, and the tensors are represented in *NHWC* format. On the left side, the filter size is (3×3) with a padding of 1; on the right side, the filer size is (5×5) with a padding of 2.

5.3. Im2col-Winograd

Winograd (Lavin & S.Gray, 2016) has been extensively used to accelerate convolutions. When using filters of size r to



calculate n outputs, Winograd $F(n, r)$ needs $(n + r - 1)$ multiplications, less than $(n * r)$ multiplications required by standard convolution, thereby improving efficiency.

Previous studies have implemented Winograd on GPUs to accelerate 2D convolution (Castro et al., 2021; Chetlur et al., 2014; Yan et al., 2020; Yang & Lai, 2021). These studies mainly use the 2D variant of Winograd, and arrange tensors in *NCHW* or *CHWN* format. The fused-Winograd2D is limited to (3×3) filters, while the non-fused has been applied to other filter sizes but requires a large amount of memory to store intermediate variables. *NHWC* is a widely used tensor format, but it is not well-suited for Winograd2D due to its discontinuous memory access, which reduces the hit ratio of GPU L2-cache.

To implement a versatile and flexible fused-Winograd on GPUs for *NHWC* format, we propose Im2col-Winograd. As shown in Figure 8, Im2col-Winograd has 2 stages: lower the filters and feature-maps to 2D matrices through Im2col; cumulatively apply 1D Winograd on these matrices. The 2 stages are integrated into a single operator, eliminating the need for auxiliary memory.



Figure 7. Im2col-Winograd with 4/8/16 states.

$F(n, r)$ uses $\alpha = (n + r - 1)$ variables to get n outputs, so it can be denoted as α state Winograd. Fused-Winograd loads, transforms, and stores the data in shared-memory, where the max shared-memory size for 1 block is 49152 bytes. To hide memory latency and maximize item-reuse, GPUs typically perform $8 * (8 \times 8)$ outer products via 256 threads within a block. To satisfy such outer products, the state α must ≤ 16 , and ideally a power of 2. $F(2 \times 2, 3 \times 3)$ for (3×3) filters has 16 states, exactly meeting the upper limit.

As shown in Figure 7, 4/8/16 state Im2col-Winograd has

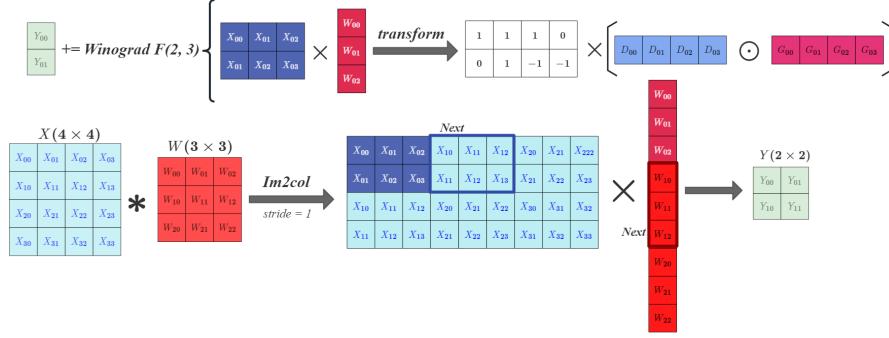


Figure 8. Img2col-Winograd $F(2, 3)$. The filter W and input-feature-map X are converted to 2D matrices through Im2col. The items selected by sliding windows are used to perform $F(2, 3)$, and the results are accumulated to the output-feature-map Y .

been implemented to cover 2-9 filter-sizes for both the forward and backward propagation. Compared to GEMM convolution, the 4/8 state Im2col-Winograd exhibits almost no loss of precision, whereas the 16-state version yields a relative difference of 10^{-4} . Both the 8 and 16 state versions can be manually enabled or disabled.

Im2col-Winograd is more flexible and lightweight than Winograd2D. Unlike Winograd2D, Im2col-Winograd only imposes constraints on one dimension of filters, as opposed to two. In some cases, Im2col-Winograd requires fewer resources than Winograd2D, to achieve the same acceleration. For instance, both $F(2 \times 2, 3 \times 3)$ and $F(6, 3)$ theoretically reduce the multiplication to $1/2.25$. However, the latter needs 8 states and fetches 33/6 items to calculate an output, while the former uses 16 states and selects 25/4 items. Moreover, Im2col-Winograd can be used for ND convolution, by extending Im2col from 2D to ND.

Im2col-Winograd has been compared with cuDNN-8.9 (Chetlur et al., 2014) on RTX3060ti GPU. As shown in Figure 9, Im2col-Winograd especially its 16-state version is more efficient than cuDNN in many cases.

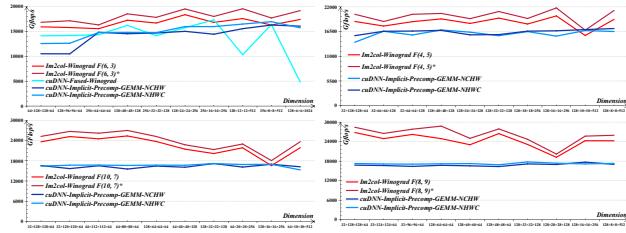


Figure 9. Compare Im2col-Winograd with cuDNN. $F(6, 3)$, $F(4, 5)$, $F(10, 7)$ and $F(8, 9)$ are performed with (3×3) , (5×5) , (7×7) and (9×9) filters respectively. The tensors are represented in NHWC format. In forward propagation, filters are transposed from NHWC to HWCN format with a small expense to improve bandwidth. '*' means ignoring the time of filter-transposition.

5.4. Parallel Random Number Generation

Cu32 generates pseudo-random numbers via multi threads. Each thread performs linear congruence on global seeds to obtain its local seeds, and then conducts linear congruence on local seeds to produce random numbers. Users can specify global seeds or have them randomly generated.

The quality of random numbers used to initialize DNNs significantly impacts the convergence. We have found 2 factors that affect such quality: the volume of random numbers generated by individual threads, and the degree of randomness introduced in local seeds. Figure 10 illustrates the influence of these 2 factors on the convergence of ResNet18 on Cifar10.

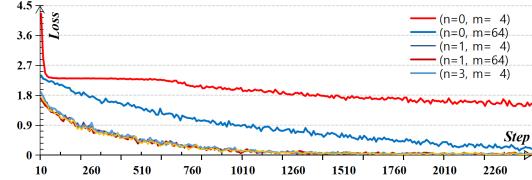


Figure 10. The quality of random numbers impacts the convergence. (n, m) means performing n times linear-congruence on global seeds, and producing m random numbers per thread

5.5. Inplace Operators and Fused Operator

In DNNs, certain operators have predetermined combination orders, and therefore can be fused or inplace to save memory and bandwidth. Furthermore, compared to DL frameworks that process the code, programmers often have more prior knowledge of the DNNs. This enables them to undertake certain optimizations on the computational graphs using inplace or fused operators. Alpha provides inplace and fused operators for users, as demonstrated in Figure 11.

```

nn.batchNorm_leakyRelu(nn.batchNorm(inplace, feature_dim, nn.leakyRelu(negative_slope));
nn.batchNorm_relu(nn.batchNorm(feature_dim, nn.relu(inplace)));
nn.global_batchNorm_leakyRelu(nn.global_batchNorm(feature_dim), nn.leakyRelu());
nn.affine_leakyRelu(nn.affine(inplace, feature_dim, nn.leakyRelu()));
nn.leakyRelu_dropout(nn.leakyRelu(inplace, negative_slope), nn.dropout(inplace, nonzero_prop));

F.add_leakyRelu(negative_slope, XI, X2);
F.log_softmax(crossEntropy(features));
loss.sigmoid.binaryCrossEntropy();

eg.intg_adjust_color(XI.to_int8(inplace), brightness, saturation, contrast);
eg.intg_affine.rotate(theta).shear(shy, shx).translate(ty, tx)
    .transform(XI, inplace, height, width);
    
```

Figure 11. The inplace and fused operators APIs of Alpha.

5.6. Accuracy Optimizations

For reduction operators, Kahan-summation is used to reduce errors. Besides, 'minimizing float operations' and 'avoiding numerical overflow' can be balanced. For example, when computing the mean of collection $\{x_1, x_2, \dots, x_n\}$, $\frac{1}{N} \sum x_i$ has better precision than $\sum \frac{1}{N} x_i$

due its fewer float operations, but it takes higher risk of numerical overflow. To compromise between these 2 methods, the collection can be divided into M segments, and the sum of each segment S_i can be calculated to determine the average $\frac{1}{N} \sum S_i$.

GELU (Hendrycks & Gimpe, 2016) is commonly estimated as $0.5(1 + \tanh[\sqrt{\pi/2}(x + 0.044715x^3)])$, with its derivative $\frac{1}{1+e^u}(1 - \frac{e^u(u-0.14271x^3)}{1+e^u})|_{u=-(1.59577x+0.214064x^3)}$. However, e^{-u} may yield *inf* leading to *NaN* derivative, when $x < 0$. To avoid this, the derivative can be calculated as $\frac{e^{-u}}{1+e^{-u}}(1 - \frac{e^u(u-0.14271x^3)}{1+e^{-u}})$.

6. Experiments

To evaluate Alpha's performance, certain comparisons with PyTorch&cuDNN have been conducted on Cifar10 and ILSVRC2012. PyTorch version is 1.12, CUDA version is 11.5, and the version of Alpha&cu32 is 1.1 or 1.2.

6.1. Methods and Conditions

Several kinds of optimizers (Rumehart et al., 1986; Tielemans & Hinton, 2012; Kingma & Ba, 2015; Loshchilov & Hutter, 2017; Liu et al., 2019) were utilized to train DNNs (Szegedy et al., 2015; Simonyan & Zisserman, 2015; He et al., 2016; Hu et al., 2018), with softmax and 0.001 learning-rate on Cifar10 and ILSVRC2012:

- Cifar10: the input shape is (32×32) , the label space is 10, the batchsize is 512, and the data was processed on an RTX3060ti GPU with an i5-12490 CPU.
- ILSVRC2012: the input shape is (128×128) , the label space is 1000, the batchsize is 256, and the data was processed on an RTX4090 with an i9-13900KF CPU. Alpha uses 16 threads to load and preprocess data, while PyTorch utilizes 4 workers, which is the optimal configuration we have observed.

DNNs of Alpha and PyTorch were identical, and underwent the same initialization, training, and testing procedures. The activation function is LeakyRelu (Maas et al., 2013). Specific convolutional and full-connect layers were adjusted to accommodate the tensor shape, while the backbone of DNNs remains unaltered. BatchNorm (Ioffe & Szegedy, 2015) was integrated into VGG and GoogLeNet, to prevent gradient-vanishing and expedite convergence. Full-connect and convolutional layers were initialized using kaiming-uniform (He et al., 2015).

The labels were encoded to one-hot formats, and the pixel values were linearly scaled to fall within $[-1, 1]$. The loss-function value was recorded every 10 steps. To plot the loss curves of ILSVRC2012, a sliding window of length 10 was used to average the loss-function values without overlap.

6.2. Results and Discussions

Table 1-2 present the performance of Alpha and PyTorch. The loss curves are shown in figure 12-35 in *Appendix*.

Table 1. Performance on Cifar10. PyTorch's data is blue. Alpha's data is red. The speed and acceleration of Alpha-1.2 is dark red.

Network	Training	Speed	Acceleration	Train/Test accuracy	GPU memory	Reduction	GPU utilization	CPU utilization	Weight file
GoogLeNet	Adam 30 epoch	9.147x 8.674 epochs	1.423x 1.501x	97.82% 79.78%	2121 <small>ms</small>	0.4491x	90% 183 W	18.0% 1067 ms	32.6 ms
	SGDM 40 epoch	9.125x 8.665 epochs	1.421x 1.496x	92.86% 63.22%	2089 <small>ms</small>	0.4437x	98% 183 W	12.8% 3814 ms	23.8 ms
	Adam 35 epoch	5.88x 5.332 epochs	1.256x 1.371x	99.69% 78.11%	1867 <small>ms</small>	0.4292x	92% 197 W	13.6% 1138 ms	66.7 ms
ResNet18	Adam 35 epoch	5.794x 5.284 epochs	1.235x 1.354x	100.0% 60.00%	1010 <small>ms</small>	0.4099x	98% 162 W	12.6% 3012 ms	48.2 ms
	SGDM 35 epoch	7.158x		99.83% 61.23%	2464 <small>ms</small>	0.4099x			
	Adam 30 epoch	11.528x 10.464 epochs	1.227x 1.352x	99.01% 79.45%	1685 <small>ms</small>	0.5397x	92% 198 W	13.6% 1188 ms	120 ms
ResNet34	SGDM 35 epoch	11.471x 10.269 epochs	1.202x 1.343x	100.0% 60.00%	1672 <small>ms</small>	0.5045x	98% 165 W	12.6% 3813 ms	87.3 ms
	Adam 35 epoch	11.388x		99.83% 61.23%	3116 <small>ms</small>	0.5045x			
	Adam 35 epoch	9.789x 8.965 epochs	1.137x 1.289x	97.91% 82.75%	1634 <small>ms</small>	0.4466x	97% 197 W	12.9% 1007 ms	78.7 ms
VGG16	SGDM 35 epoch	9.698x 8.993 epochs	1.121x 1.279x	100.0% 75.99%	1569 <small>ms</small>	0.4285x	99% 180 W	12.7% 3114 ms	56.7 ms
	Adam 35 epoch	10.871x		100.0% 75.99%	3658 <small>ms</small>	0.4289x			
	Adam 40 epoch	11.849x 10.389 epochs	1.109x 1.265x	96.00% 81.13%	1786 <small>ms</small>	0.4756x	97% 197 W	12.7% 1042 ms	106 ms
VGG19	SGDM 40 epoch	13.142x		96.03% 80.98%	3755 <small>ms</small>	0.4529x	99% 181 W	12.7% 3817 ms	77 ms
	Adam 40 epoch	11.802x 10.268 epochs	1.098x 1.262x	99.94% 76.84%	1694 <small>ms</small>	0.4529x			
	SGDM 40 epoch	12.364x		99.60% 76.58%	3740 <small>ms</small>	0.4529x			
SENet	Adam 30 epoch	25.767x 25.288 epochs	1.138x 1.159x	99.63% 81.88%	5424 <small>MB</small>	0.77808x	94% 197 W	13.1% 1301 ms	207 ms
	SGDM 30 epoch	29.317x		98.66% 81.46%	6947 <small>MB</small>	0.77808x			
	SGDM 35 epoch	25.718x 25.267 epochs	1.121x 1.141x	100.0% 63.74%	5261 <small>ms</small>	0.7594x	99% 187 W	12.3% 3116 ms	150 ms

Table 2. Performance on ILSVRC2012. PyTorch's data is blue.

Alpha's data is red. The speed and acceleration of Alpha-1.2 is dark red.

Network	Training	Speed	Acceleration	Train set accuracy	GPU memory	Reduction	GPU utilization	CPU utilization	Weight file
ResNet18	Adam 50 epoch	664.957x 627.675 epochs	1.809x 1.916x	99.99% 98.90%	5718 <small>ms</small>	0.4826x	92% 388 W	61.9% 5360 <small>mb</small>	66.8 ms
	RMsgprop 40 epoch	661.992x 625.202 epochs	1.790x 1.896x	97.78% 97.53%	5656 <small>ms</small>	0.4742x	98% 368 W	24.7% 7321 <small>mb</small>	50.9 ms
	Adam 60 epoch	1264.301x 1122.489 epochs	1.756x 1.906x	99.01% 98.95%	9090 <small>ms</small>	0.5490x	95% 405 W	36.5% 5180 <small>mb</small>	124 ms
ResNet34	SGDM 60 epoch	2219.791x 2126.331 epochs	1.788x 1.958x	98.77% 97.99%	15384 <small>ms</small>	0.5979x	99% 370 W	23.8% 7252 <small>mb</small>	39.8 ms
	Adam 60 epoch	1233.389x 1126.331 epochs	1.756x 1.906x	99.04% 98.95%	10538 <small>ms</small>	0.5979x			
	AdamW 60 epoch	1208.236x 1126.236 epochs	1.788x 1.958x	98.85% 97.53%	12896 <small>ms</small>	0.7733x	91% 389 W	45.1% 5406 <small>mb</small>	264 ms
ResNet50	AdamW 60 epoch	989.484x 975.828 epochs	1.489x 1.509x	98.85% 98.30%	9980 <small>ms</small>	0.7733x	99% 371 W	20.6% 7357 <small>mb</small>	199 ms
	SGDM 60 epoch	1472.767x 1462.853 epochs	1.063x 1.168x	98.63% 98.21%	12896 <small>ms</small>	0.7733x			
	Adam 50 epoch	1088.315x 932.007 epochs	1.025x 1.196x	97.94% 97.29%	10870 <small>ms</small>	0.7633x	96% 402 W	40.3% 5106 <small>mb</small>	294 ms
VGG16	SGDM 40 epoch	1114.341x 1042.222 epochs	1.024x 1.196x	94.81% 94.01%	13564 <small>ms</small>	0.7854x	97% 366 W	23.3% 7294 <small>mb</small>	224 ms
	Adam 40 epoch	1217.609x 1023.881 epochs	1.024x 1.207x	97.45% 97.27%	11138 <small>ms</small>	0.7711x	98% 405 W	38.7% 5031 <small>mb</small>	319 ms
	SGDM 40 epoch	1246.722x 1129.833 epochs	1.016x 1.209x	97.30% 97.35%	10880 <small>ms</small>	0.7903x	97% 370 W	23.1% 7460 <small>mb</small>	244 ms
VGG19	Adam 60 epoch	1209.301x 1016.869 epochs	1.016x 1.209x	96.98% 96.88%	10268 <small>ms</small>	0.8384x	90% 355 W	46.7% 5627 <small>mb</small>	208 ms
	SGDM 60 epoch	1346.629x 1032.983 epochs	1.480x 1.491x	97.30% 96.25%	12242 <small>ms</small>	0.8384x	94% 357 W	21.4% 7301 <small>mb</small>	152 ms
	Adam 60 epoch	916.739x 1033.822 epochs	1.457x 1.479x	97.30% 96.25%	10268 <small>ms</small>	0.8384x			

On both datasets, DNNs trained by Alpha or PyTorch show comparable convergence, accuracy, and efficiency. In terms of speed and memory-usage, Alpha outperforms PyTorch in certain scenarios.

Several factors contribute to Alpha's higher speed compared to PyTorch. Cu32 provides efficient GPU operators, while ensuring correctness and accuracy. The C-K-S algorithm avoids redundant 0-calculations in conv-layers, reducing time complexity and strain on hardware. Specifically, KS-deconv and Sk-dilated significantly simplify the backpropagation of conv-layers with non-unit stride, thereby accelerating the execution of ResNet and SENet. Im2col-Winograd reduces the cost of stride-1 conv-layers, especially those with (3×3) or (5×5) filters. The concurrent execution offered by async APIs enhances the power of GPUs, and conceals the execution of certain CPU functions. The fusion of operators reduces memory access, lowering the computational cost.

Alpha's lower memory-usage compared to PyTorch can be attributed to several reasons. Alpha offers a wider range of inplace and fused operators than PyTorch; and it tries to calculate gradients directly within the memory space of preceding gradients, thereby reducing the need for memory space to store intermediate variables. In both forward and backward propagation, Tensors that are no longer required can be released based on Syncers, thus facilitating the

recycling of memory blocks. The experiments mentioned in this paper were carried out using the basic memory pool, and it is possible to further decrease memory usage by configuring a more advanced memory-pool to EngineCore. GPU-utilization is evaluated using 2 metrics: power and time-slice utilization-rate (UT). Alpha and PyTorch employ distinct methods to scheduling operators, each offering its own advantages. In PyTorch, operators are launched by a single thread, and typically executed in 1 CUDA stream based on a first-in-first-out order. This approach leads to higher UT due to less overhead of synchronization between operators. Conversely, in Alpha, operators are managed via multi-threads, and executed in multiple CUDA streams concurrently, enabling higher power. However, this approach requires heavier synchronization mechanisms, leading to a lower UT.

CPU-utilization is evaluated by 2 metrics: memory-usage and time-slice utilization-rate (UT). The memory-usage is primarily used to indicate the absence of memory leaks, as it can be influenced by the configuration of the Java or Python virtual machine. When training DNNs on ILSVRC2012, Alpha and PyTorch achieve nearly 100% GPU utilization, suggesting that the speed bottleneck is GPU-computing rather than the data loading and preprocessing on CPU. Notably, Alpha's 16 threads have higher parallelism than PyTorch's 4 workers, and therefore result in a higher UT. Overall, these experimental results are evidence supporting the correctness and efficiency of Alpha&cu32.

7. Conclusion

This paper introduces the background, characteristics, architectures, and techniques of Alpha&cu32. Besides, it demonstrates their correctness and effectiveness, by benchmarking them against PyTorch&cuDNN on Cifar10 and ILSVRC2012.

Alpha tries to raise Java's effectiveness in DL field. It adopts some innovative designs to ensure user-friendliness, and realizes high-performance through cu32 and specific techniques. Based on its multi-layer architecture and the adaptability to Java big-data ecosystem, Alpha achieves scalability and is capable of aggregating heterogeneous computing resources.

The source code for Alpha is available for access at {the link is hidden due to double-blind review}. Subsequent updates will be aligned with our future works.

Statement

This paper presents work whose goal is to advance the field of Machine Learning. There are many potential societal consequences of our work, none which we feel must be specifically highlighted here

References

- Abadi, M., Agarwal, A., Barham, P., and et al. TensorFlow: Large-scale machine Learning on heterogeneous system. *arXiv e-prints*, 2016. URL <https://arxiv.org/abs/1603.04467v2>.
- Castro, R. L., Andrade, D., and Fraguela, B. B. OpenCNN: A Winograd Minimal Filtering Algorithm Implementation in CUDA. *Mathematics*, 9(17), 2021.
- Chetlur, S., Woolley, C., Vandermersc, P., and et. al. cuDNN: Efficient Primitives for Deep Learning. *CoRR*, abs/1410.0759, 2014, 2014.
- Chollet, F. Keras: Deep Learning for humans. URL <https://keras.i>.
- Cutting, D., Cafarella, M., and et. al. Hadoop. URL <https://hadoop.apache.org/>.
- Ding, J., Ren, X., Luo, R., and Sun, X. An Adaptive and Momental Bound Method for Stochastic Learning. *arXiv e-prints*, 2019. URL <https://arxiv.org/abs/1910.12249v1>.
- Feenster, A., Henkelmann, C., Bamberg, E., and et. al. Deep Java Library. URL <https://djl.ai/>.
- Gibson, A., Nicholson, C., Patterson, J., and et al. Deeplearning4j: Distributed, open-source deep learning for Java and Scala on Hadoop and Spark. 2016. doi: 10.6084/M9.FIGSHARE.3362644.
- Goodfellow, I., Bengio, Y., and Courville, A. *Deep Learning*. MIT Press, 2016. URL <http://www.deeplearningbook.org>.
- He, K., Zhang, X., Ren, S., and Sun, J. Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet. In *Proceedings of International Conference on Machine Learning (ICML)*, pp. 448–456, 2015.
- He, K., Zhang, X., Ren, S., and Sunn, J. Deep residual learning for image recognition. In *Proceedings of IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 770–778, 2016. doi: 10.1109/CVPR.2016.90.
- Hendrycks, D. and Gimpe, K. Gaussian Error Linear Units (GELUs). *arXiv e-prints*, 2016. URL <https://arxiv.org/abs/1606.08415v5>.
- Hu, J., Shen, L., and Sun, G. Squeeze-and-Excitation Networks. In *Proceedings of IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 7132–7141, 2018. doi: 10.1109/CVPR.2018.00745.

- 440 Ioffe, S. and Szegedy, C. Batch Normalization: Accelerating
 441 Deep Network Training by Reducing Internal Covariate
 442 Shift. In *Proceedings of International Conference on*
 443 *Machine Learning (ICML)*, pp. 448–456, 2015.
- 444 Jia, Y., Shelhamer, E., Donahue, J., and et al. caffe:
 445 Convolutional architecture for fast feature embedding. In
 446 *Proceedings of the 22nd ACM international conference*
 447 *on Multimedia Retrieval (ICMR)*, pp. 675–678, Newark,
 448 NJ, USA, 2014.
- 449 Kingma, D. P. and Ba, J. L. Adam: a Method for Stochastic
 450 Optimization. In *Proceedings of the 3rd International*
 451 *Conference on Learning Representations (ICLR)*, San
 452 Diego, USA, 2015.
- 453 Krizhevsky, A. Cifar10, a. URL <http://www.cs.toronto.edu/~kriz/cifar.html>.
- 454 Krizhevsky, A. cuda-convnet2, b. URL <https://code.google.com/p/cuda-convnet2>.
- 455 Krizhevsky, A., Sutskever, I., and Hinton, G. E. Imagenet
 456 classification with deep convolutional neural networks.
 457 In *Proceedings of the 26th Annual Conference on Neural*
 458 *Information Processing Systems (NIPS)*, pp. 1097–1105,
 459 2012.
- 460 Lavin, A. and S. Gray. Fast algorithms for convolutional
 461 neural networks. In *Proceedings of IEEE Conference on*
 462 *Computer Vision and Pattern Recognition (CVPR)*, Las
 463 Vegas, NV, USA, 2016. doi: 10.1109/CVPR.2016.435.
- 464 Li, F. IMAGENET. URL <https://image-net.org>.
- 465 Liu, L., Jiang, H., He, P., and et. al. On the Variance
 466 of the Adaptive Learning Rate and Beyond. *arXiv e-prints*,
 467 2019. URL <https://arxiv.org/abs/1908.03265v1>.
- 468 Loshchilov, H. and Hutter, F. Decoupled Weight decay
 469 regularization. *arXiv e-prints*, 2017. URL <https://arxiv.org/abs/1711.05101v1>.
- 470 Maas, A. L., Y. A., Hannun, and Ng, A. Y. Rectifier
 471 Nonlinearities Improve Neural Network Acoustic Models.
 472 In *Proceedings of International Conference on Machine*
 473 *Learning (ICML)*, 2013.
- 474 NVIDIA. cuBLAS. URL <https://developer.nvidia.cn/cublas>.
- 475 NVIDIA. *CUDA C Best Practices Guide*. NVIDIA,
 476 2023a. URL <https://docs.nvidia.com/cuda/cuda-c-best-practices-guide>.
- 477 NVIDIA. *CUDA C Programming Guide*. NVIDIA,
 478 2023b. URL <https://docs.nvidia.com/cuda/cuda-c-programming-guide>.
- 479 Paszke, A., Gross, S., Massa, and et.al. Spark: Cluster
 480 Computing with Working Sets. *USENIX Association*,
 481 2010.
- 482 Paszke, A., Gross, S., Massa, F., and et al. PyTorch:
 483 An Imperative Style, High-Performance Deep Learning
 484 Library. In *Proceedings of the 33rd Annual Conference*
 485 *on Neural Information Processing Systems (NIPS)*, pp.
 486 675–678, Vancouver, Canada, 2019.
- 487 Rumehart, D. E., Hinton, G. E., and et al. Learning
 488 Representations by Back-Propagating errors. *Nature*, 323
 489 (6088):533–536, 1986. doi: 10.1038/323533a.
- 490 Simonyan, K. and Zisserman, A. Very Deep Convolutional
 491 Networks for Large-Scale Image Recognition. In
 492 *Proceedings of the 3rd International Conference on*
 493 *Learning Representations (ICLR)*, 2015.
- 494 Srivastava, N., Hinton, G. E., Krizhevsky, A., and et al.
 495 Dropout: A Simple Way to Prevent Neural Network from
 496 Overfitting. *Journal of Machine Learning Research*, 15
 497 (1):1929–1958, 2014.
- 498 Szegedy, C., Liu, W., Jia, Y., and et al. Going deeper with
 499 convolutions. In *Proceedings of IEEE Conference on*
 500 *Computer Vision and Pattern Recognition (CVPR)*, pp.
 501 1–9, 2015. doi: 10.1109/CVPR.2015.7298594.
- 502 Tieleman, T. and Hinton, G. RMSprop. *Neural Networks*
 503 for Machine Learning, 2012.
- 504 Tokui, S., Oono, K., Hido, S., and Clayton, J. Chainer: a
 505 next-generation open-source framework for deep learning.
 506 In *Proceedings of the 29th Annual Conference on*
 507 *Neural Information Processing Systems (NIPS)*, Montréal
 508 Canada, 2015.
- 509 Yan, D., Wang, W., and Chu, X. Optimizing Batched
 510 Winograd Convolution on GPUs. In *Proceedings of*
 511 *the 25th ACM Symposium on Principles and Practice*
 512 *of Parallel Programming (SIGPLAN)*, San Diego, CA,
 513 USA, 2020.
- 514 Yang, J. L. D. and Lai, J. Optimizing Winograd-Based
 515 Convolution with Tensor Cores. In *Proceedings of the*
 516 *50th International Conference on Parallel Processing*,
 517 2021.

Appendix

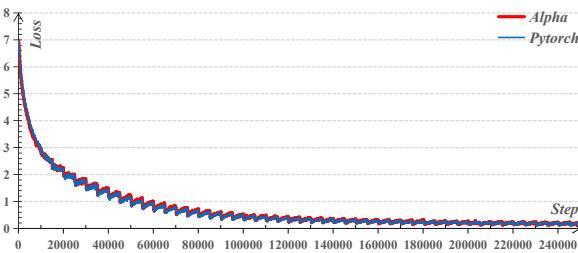


Figure 12. ILSVRC2012: ResNet18 + Adam, 50 epoch

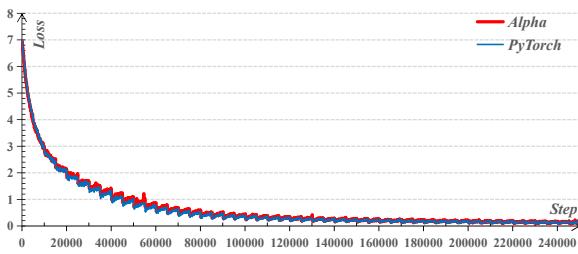


Figure 14. ILSVRC2012: ResNet34 + Adam, 50 epoch

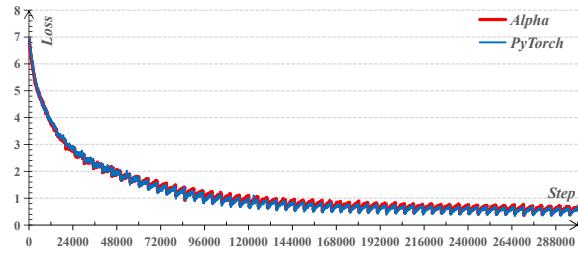


Figure 16. ILSVRC2012: ResNet50 + AdamW, 60 epoch

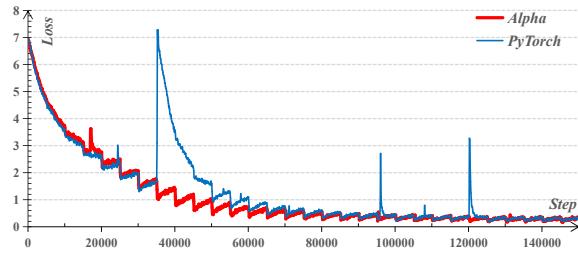


Figure 18. ILSVRC2012: VGG16 + Adam, 30 epoch

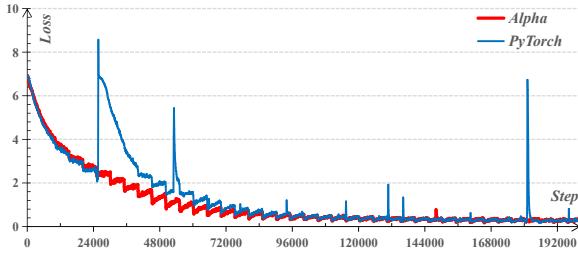


Figure 20. ILSVRC2012: VGG19 + Adam, 40 epoch

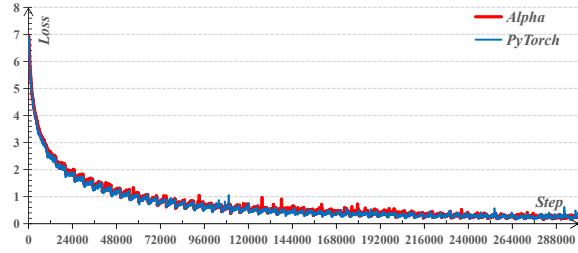


Figure 22. ILSVRC2012: SENet + RAdam, 60 epoch

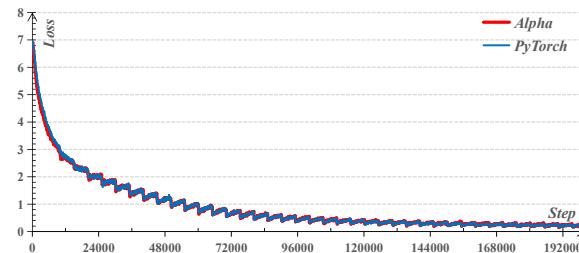


Figure 13. ILSVRC2012: ResNet18 + RMSprop, 40 epoch

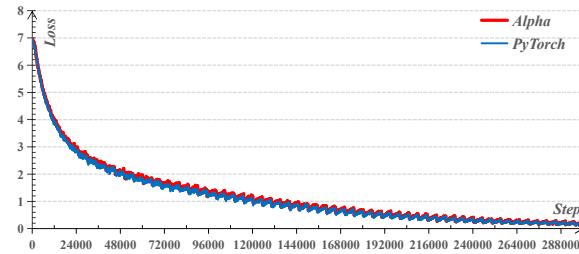


Figure 15. ILSVRC2012: ResNet34 + SGDM, 60 epoch

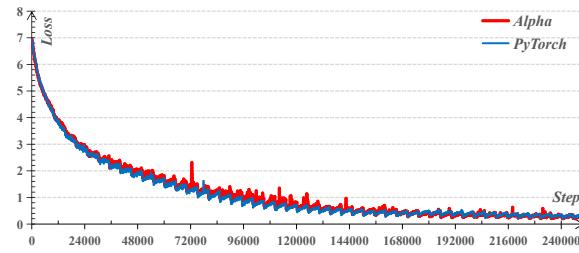


Figure 17. ILSVRC2012: ResNet50 + Adam, 50 epoch

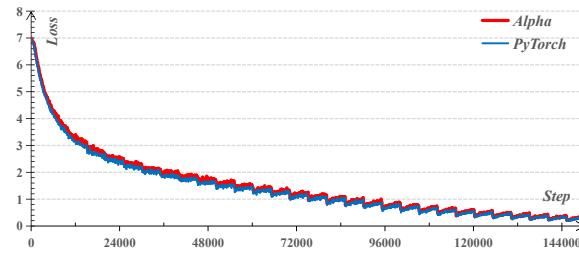


Figure 19. ILSVRC2012: VGG16 + SGDM, 30 epoch

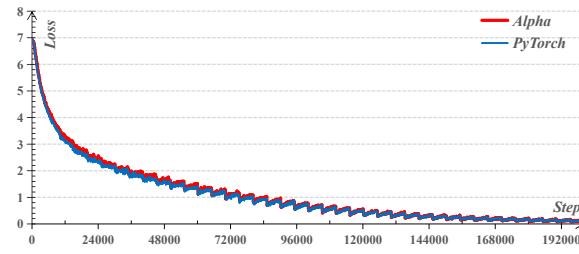


Figure 21. ILSVRC2012: VGG19 + SGDM, 40 epoch

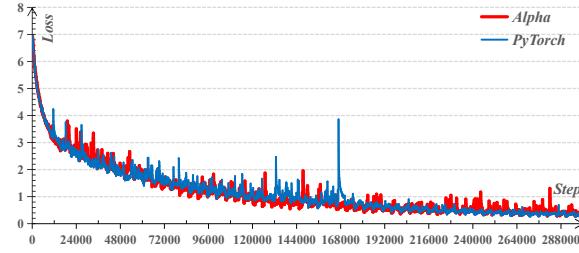


Figure 23. ILSVRC2012: SENet + Adam, 60 epoch

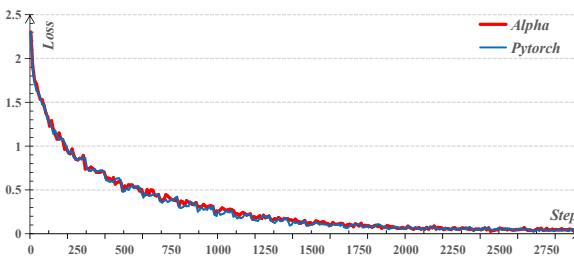


Figure 24. Cifar10: GoogLeNet + Adam, 30 epochs

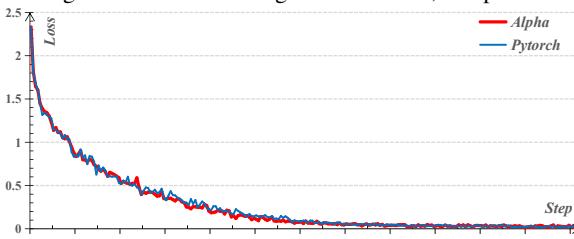


Figure 26. Cifar10: ResNet18 + Adam, 25 epoch

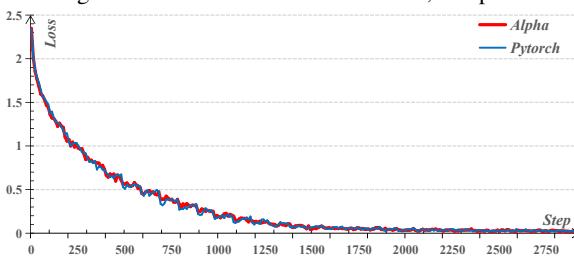


Figure 28. Cifar10: ResNet34 + Adam, 30 epoch

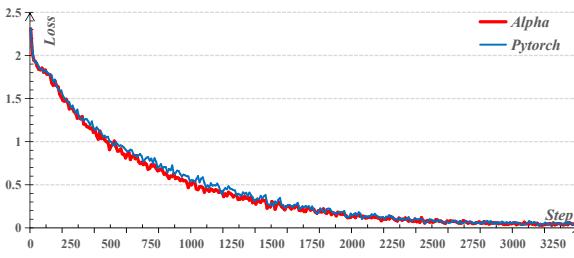


Figure 30. Cifar10: VGG16 + Adam, 35 epoch

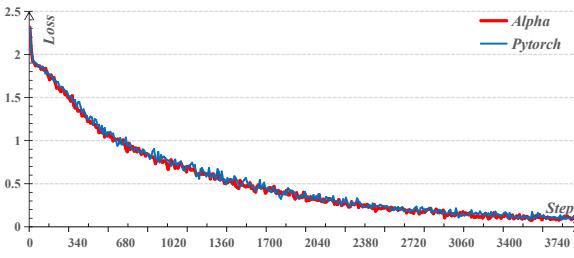


Figure 32. Cifar10: VGG19 + Adam, 40 epoch

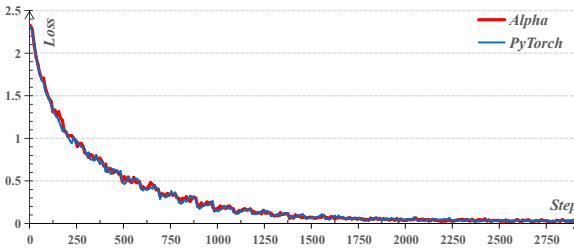


Figure 34. Cifar10: SENet + Adam, 30 epoch

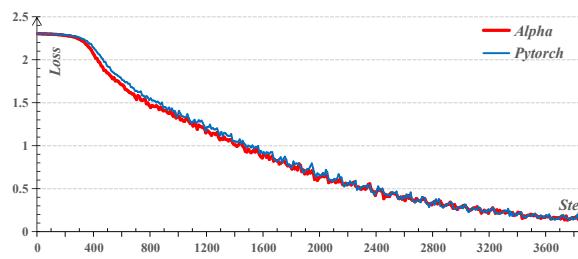


Figure 25. Cifar10: GoogLeNet + SGDM, 40 epoch

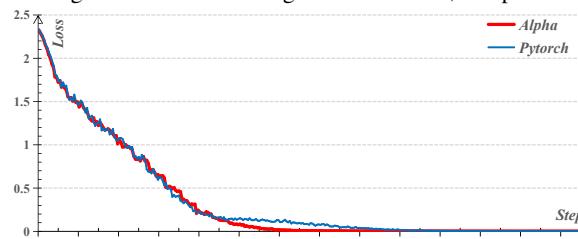


Figure 27. Cifar10: ResNet18 + SGDM, 35 epoch

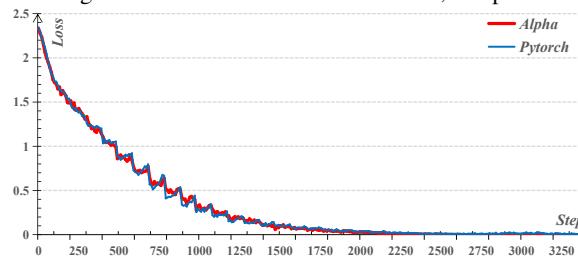


Figure 29. Cifar10: ResNet34 + SGDM, 35 epoch

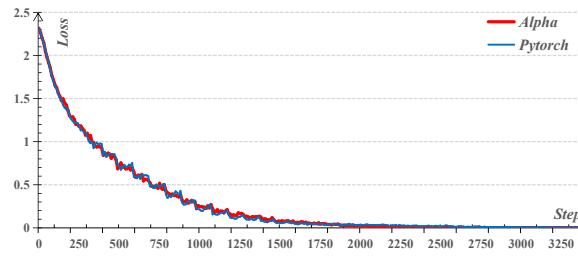


Figure 31. Cifar10: VGG16 + SGDM, 35 epoch

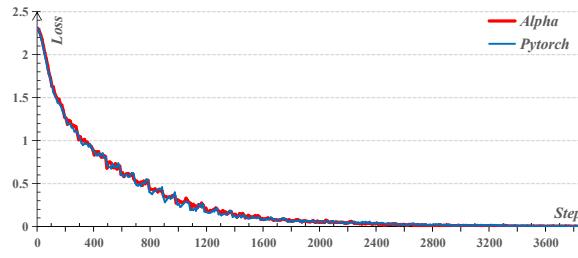


Figure 33. Cifar10: VGG19 + SGDM, 40 epoch

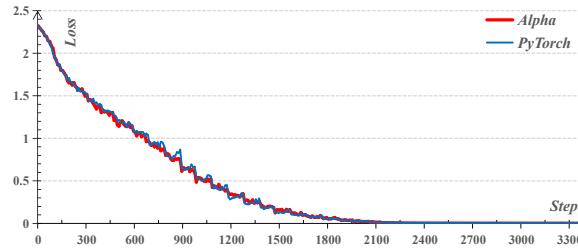


Figure 35. Cifar10: SENet + SGDM, 35 epoch